

TP1F3B

July 18, 2017

1 Introduction to data mining with the tips dataset

- Romain Billot
- Yannis Haralambous
- Philippe Lenca
- Sorin Moga

1.1 Data and objective understanding

–The Tips dataset– Food server’s tips in restaurants may be influenced by many factors (e.g. the nature and location of the restaurant, the size of the party, the table location and the day of the week. . .). Restaurant managers need to know which factors matter when they assign tables to food servers. Indeed, for the sake of staff morale, they usually want to avoid either the substance or the appearance of unfair treatment of the servers, for whom tips (at least in restaurants in the United States) are a major component of pay. In one restaurant, a food server recorded some data on all customers they served during an interval of two and a half months in early 1990. The restaurant, located in a suburban shopping mall, was part of a national chain and served a varied menu. In observance of local law the restaurant offered seating in a non-smoking section to patrons who requested it. Each record includes a day and time, and thus taken together, they show the server’s work schedule. The food server provided a comma-separated-value file tips.csv containing 244 records, described by 7 variables (total bill, tip, sex, smoker, day, time and size).

```
In [1]: # import useful libraries
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import numpy as np
import scipy as sp
```

1.1.1 Question 1

What do you know from the text above and what information is missing?

1.1.2 Question 2

Do you have some idea about the objectives of the study and the knowledge you could extract from the data? Could you suggest a list of questions of interest?

Here we want to know the influence of total bill, sex, smoker, day, time and size on tip by finding a model that can generalise the data.

Some questions :

- What factors can most influence the tip value ?
- Which factors matter when managers assign tables to food servers ?

1.1.3 Question 3

Load the dataset and have a look at it using the describe() function. Describe the data (the format of the data, the quantity of data –number of example/ records and variable/fields–). What are the expected values and role of each variable?

```
In [2]: # the code below loads the data for you
data_tips=pd.read_csv('tips.csv')
# you can run the below code to see the first five observations
data_tips.head()
```

```
Out[2]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
In [3]: # this function should return a tuple of 3 values: the class of your dataset
def answer_3():
    # write your code here:
    # the type of the object data_tips
    dataset_type=type(data_tips)
    # the shape of data
    shape=data_tips.shape
    #variables in data
    col_names=data_tips.columns

    return dataset_type,shape,col_names
```

```
answer_3()
```

```
Out[3]: (pandas.core.frame.DataFrame,
(244, 7),
Index(['total_bill', 'tip', 'sex', 'smoker', 'day', 'time', 'size'], dtype=object))
```

```
In [4]: # a description for each variable :
# 'total_bill' : total to pay
# 'tip' : the tip value
# 'sex' : the gender of the customer
# 'smoker' : customer smokes or not
# 'day' : 4 days ( from thursday to sunday)
# 'time' : either dinner or lunch
# 'size' : the size of the party
```

1.1.4 Question 4

Tip is usually referred to by percentage points, or as a rate. This enables a normalization over the total bill and a comparison of values across other variables. The question is now to create a “tip rate” variable and to add it to the original dataset.

```
In [5]: #This function should return a dataframe including the new variable tip_rate
def answer_4():
    # write your code here
    data_tips['tip_rate']=data_tips['tip']/data_tips['total_bill']

    return data_tips
answer_4().head()
```

```
Out[5]:
```

	total_bill	tip	sex	smoker	day	time	size	tip_rate
0	16.99	1.01	Female	No	Sun	Dinner	2	0.059447
1	10.34	1.66	Male	No	Sun	Dinner	3	0.160542
2	21.01	3.50	Male	No	Sun	Dinner	3	0.166587
3	23.68	3.31	Male	No	Sun	Dinner	2	0.139780
4	24.59	3.61	Female	No	Sun	Dinner	4	0.146808

Home work Explore the notion of scale of measurement. Provide a short note with meaningful definitions and examples. Explain why it is important to consider the right scale for each variable. What is the scale for each of the eight variables?

you can visit the link below to explore the notion of scale measurement:
<http://stattrek.com/statistics/measurement-scales.aspx?Tutorial=AP>

1.2 Descriptive statistics and visualisation

1.2.1 Question 5

Explore univariate summaries with the describe() function.

```
In [6]: #This function should return summary about numerical and categorical features
# use the output of answer4() for your this question
def answer_5():
    # write your code here:
    num_des=answer_4().describe()
    cate_des=answer_4().describe(include=['object'])
    return num_des,cate_des
answer_5()
```

```
Out[6]:
```

	total_bill	tip	size	tip_rate
count	244.000000	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672	0.160803
std	8.902412	1.383638	0.951100	0.061072
min	3.070000	1.000000	1.000000	0.035638
25%	13.347500	2.000000	2.000000	0.129127
50%	17.795000	2.900000	2.000000	0.154770


75%	24.127500	3.562500	3.000000	0.191475
max	50.810000	10.000000	6.000000	0.710345,
	sex	smoker	day	time
count	244	244	244	244
unique	2	2	4	2
top	Male	No	Sat	Dinner
freq	157	151	87	176)

1.2.2 Question 6

Plot a representation of the days distribution in the dataset and comment.

```
In [7]: # this function should return a plot of the days distribution, give a comment
def answer_6():
    # write your code here
    return sns.countplot(x='day', data=data_tips)
answer_6()
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7fed0493b828>
```




let your comment here:

1.2.3 Question 7

Prepare a plot of the amount of tips against the total bill. What can you see ? Test the correlation between the two variables.

```
In [8]: # this function should return a plot of tips against the total_bill
def answer_7():
    # write your code here
    ax = sns.regplot(x="total_bill", y="tip", data=data_tips)
    return ax
answer_7()
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7fed048ff550>
```



We can say that tip correlate strongly with the total_bill. As shown above we can draw a line that can approximate the relationship between them.

1.2.4 Question 8

Draw and interpret three boxplots :

1. the distribution of the total bill,
2. the distribution of tips;
3. the distributions of tips vs. days.

```


In [9]: def answer_8():
        # write your code here
        plt.figure()
        boxplot_total_bill = sns.boxplot(x= 'total_bill', data=data_tips)
        plt.figure()
        boxplot_tip = sns.boxplot(x= 'tip', data=data_tips)
        plt.figure()
        boxplot_tip_vs_day = sns.boxplot(x= 'tip', y='day', data=data_tips)
        return boxplot_total_bill, boxplot_tip, boxplot_tip_vs_day
answer_8()



```

```

Out[9]: (<matplotlib.axes._subplots.AxesSubplot at 0x7fed0490f710>,
        <matplotlib.axes._subplots.AxesSubplot at 0x7fecfef6e7f0>,
        <matplotlib.axes._subplots.AxesSubplot at 0x7fecfef47e10>)

```






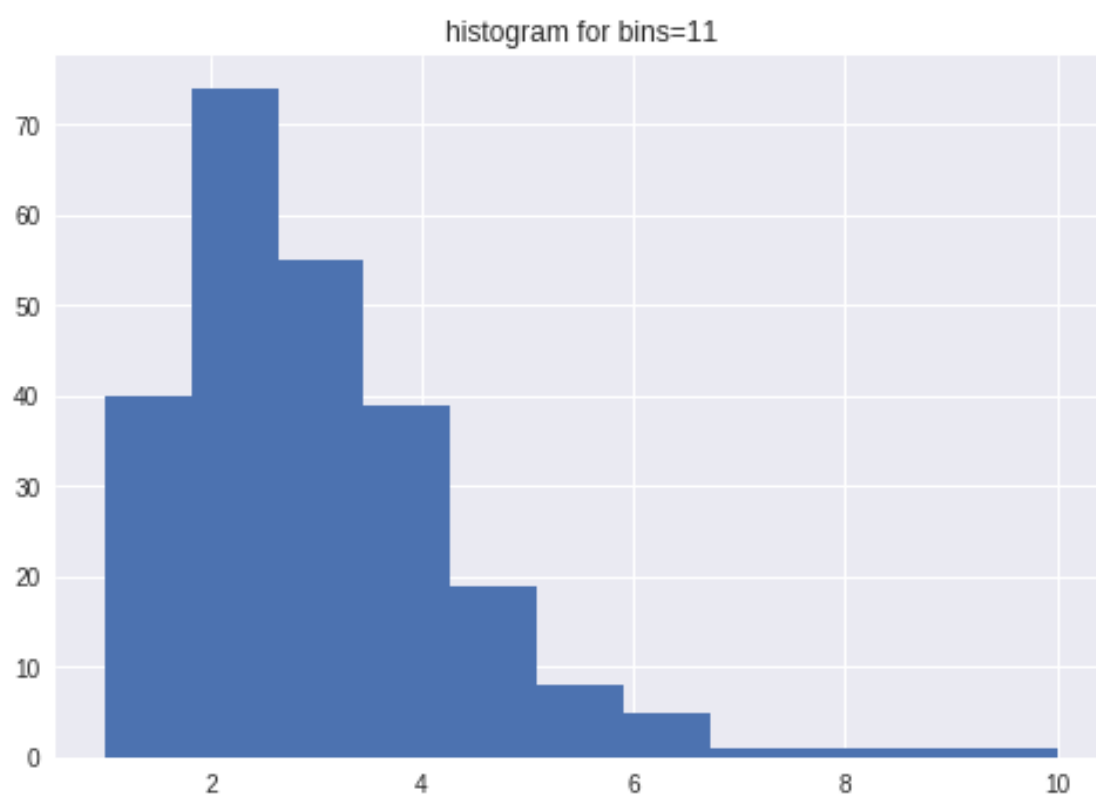
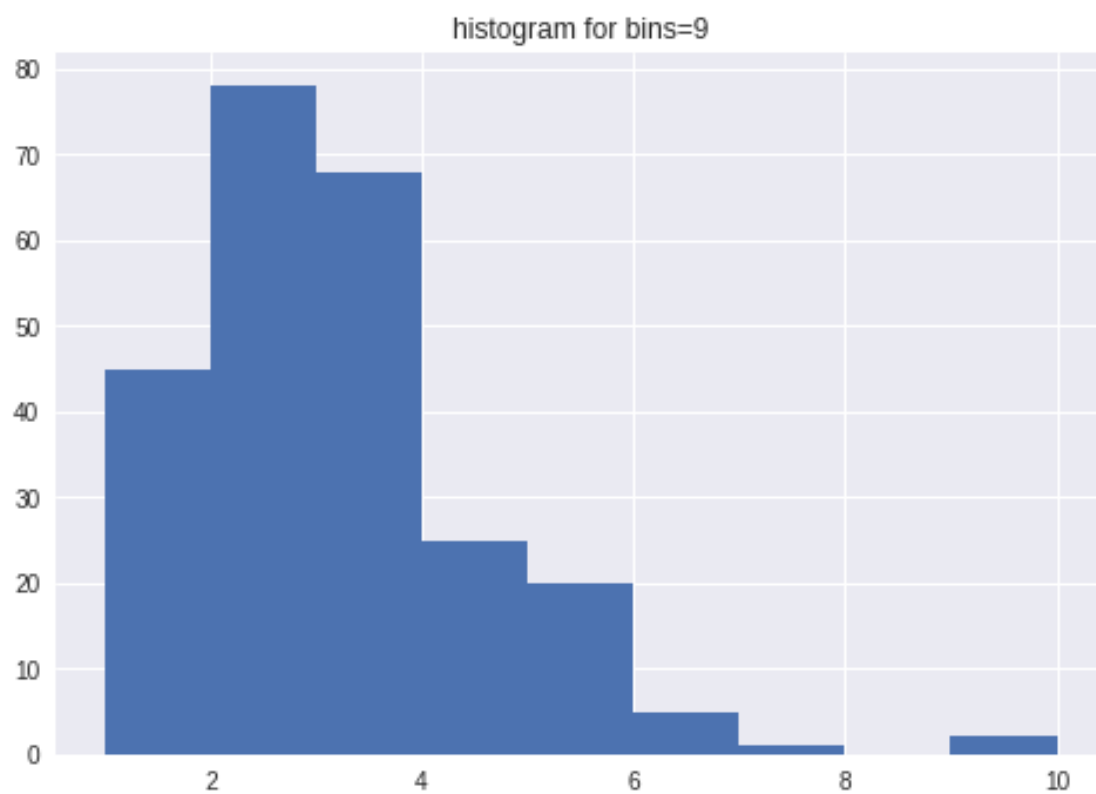
```
In [10]: # to interpret: firstly, there outliers ( black points in the boxplots) .
```


1.2.5 Question 9

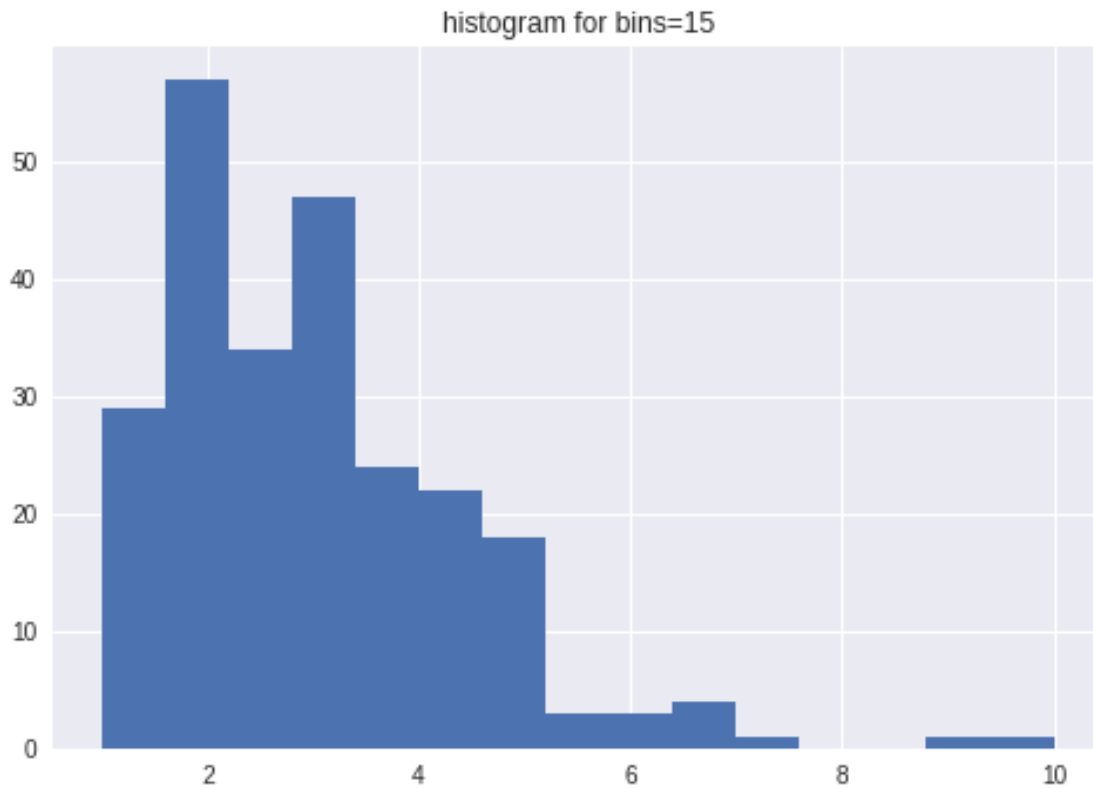
Draw an histogram of tips. What can you say about the shape of the data ? Is this restaurant expensive ? plot 6 histograms with increasing numbers of breaks.

```
In [11]: def answer_9():  
    # write your code here  
    #plt.subplot(nrows=2,ncols=3)  
    for i in range(7,17,2):  
        plt.figure()  
        plt.hist(x='tip',data=data_tips,bins=i)  
        plt.title('histogram for bins='+str(i))  
answer_9()
```











1.2.6 Question 11

Display the counts (proportions) for Gender of the Bill Payer and Smoking Parties. Do the same for time of the day (dinner or lunch) and day of the week

```
In [12]: def answer_11():
          # write your code here
          ax1=pd.crosstab(data_tips.smoker,data_tips.sex).plot.barh(stacked=True)
          ax2=pd.crosstab(data_tips.day,data_tips.time).plot.barh(stacked=True)

          return ax1,ax2
          answer_11()
```

```
Out[12]: (<matplotlib.axes._subplots.AxesSubplot at 0x7fecfec64208>,
          <matplotlib.axes._subplots.AxesSubplot at 0x7fecfeb28f98>)
```



1.2.7 Question 12

Who pay mostly the bills ? men or women ? and when ? Try to visualise the conditional distributions of Sex given the day of the week, with a mosaic plot


```
In [13]: from statsmodels.graphics.mosaicplot import mosaic
def answer_12():
    ax=mosaic(data_tips,['day','sex'])

    return ax

answer_12()
# men pay mostly and on Sunday

Out[13]: (<matplotlib.figure.Figure at 0x7fed048d2048>,
OrderedDict([('Sun', 'Female'),
(0.0, 0.0, 0.3068723249616409, 0.2360552544151075)),
('Sun', 'Male'),
(0.0,
0.23937751355132011,
0.3068723249616409,
0.7606224864486798)),
('Sat', 'Female'),
(0.3117984333360252,
0.0,
0.35128805620608894,
0.32076984763432237)),
('Sat', 'Male'),
(0.3117984333360252,
0.32409210677053502,
0.35128805620608894,
0.67590789322946487)),
('Thur', 'Female'),
(0.6680125979164984,
0.0,
0.250343212468707,
0.51441431786518055)),
('Thur', 'Male'),
(0.6680125979164984,
0.51773657700139308,
0.250343212468707,
0.48226342299860675)),
('Fri', 'Female'),
(0.9232819187595896,
0.0,
0.07671808124041025,
0.47211050883021499)),
('Fri', 'Male'),
(0.9232819187595896,
```


```
0.47543276796642764,
0.07671808124041025,
0.52456723203357236) ) ) )
```



```
In [33]: #You can explore more nice plots with the seaborn library: http://seaborn.pydata.org/
```


```
In [16]: sns.factorplot(x="day", y="tip", hue="smoker", data=data_tips) # estimotor
```

```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x7fecfe982780>
```



```
In [18]: sns.factorplot(x="day", y="total_bill", hue="sex", data=data_tips, kind="bar")
```

```
Out[18]: <seaborn.axisgrid.FacetGrid at 0x7fecfece71d0>
```



```
In [20]: # the plot above shows that Male pay mostly the bills, especially on Sunday
```

1.3 Regression

```
In [22]: # importing statsmodels
import statsmodels.formula.api as sm
```

1.3.1 Question 13

Before starting with the regression, we will learn how to build dummy variables, which is sometimes useful. Create four new variables, named day_Thur, day_Fri, day_Sat, day_Sun, that take 1 if the dining party was held on that day, 0 otherwise.

```
In [27]: #this function should return a dataframe including the dummies for 'day'
def answer_13():
    # write your code here:
    data_tips=answer_4()
    return pd.get_dummies(data_tips,columns=['day'])
answer_13().head(2)
```

```
Out[27]:
```

	total_bill	tip	sex	smoker	time	size	tip_rate	day_Fri	day_Sat	day_Sun	day_Thur
0	16.99	1.01	Female	No	Dinner	2	0.059447	0	0	0	0
1	10.34	1.66	Male	No	Dinner	3	0.160542	0	0	0	0

	day_Sun	day_Thur
0	1	0
1	1	0

1.3.2 Question 14

Fit a general linear model with tip rate as a response variable against all the other variables of interest : sex, smoker, time, size, day_Thur, day_Fri, day_Sat, day_Sun

```
In [30]: # this function should fit a general model respecting the requirements and
#( use : your_model.summary())
def answer_14():
    answer_13()
    model=sm.ols(formula="answer_13().tip_rate~answer_13().sex + answer_13().smoker + answer_13().time + answer_13().size + answer_13().day_Fri + answer_13().day_Sat + answer_13().day_Sun + answer_13().day_Thur",data=answer_13())
    return model.summary()
answer_14()
```

```
Out[30]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  answer_13().tip_rate    R-squared: 0.751
```



```

Model:                                OLS      Adj. R-squared:
Method:                            Least Squares      F-statistic:
Date:                            Tue, 18 Jul 2017      Prob (F-statistic):
Time:                            13:44:42      Log-Likelihood:
No. Observations:                244      AIC:
Df Residuals:                    236      BIC:
Df Model:                        7
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t
Intercept	0.1459	0.013	11.554	0.000
answer_13().sex[T.Male]	-0.0085	0.008	-1.023	0.307
answer_13().smoker[T.Yes]	0.0036	0.008	0.428	0.669
answer_13().time[T.Lunch]	0.0234	0.026	0.895	0.372
answer_13()['size']	-0.0096	0.004	-2.282	0.023
answer_13().day_Fri	0.0373	0.012	3.049	0.003
answer_13().day_Thur	0.0191	0.020	0.953	0.341
answer_13().day_Sun	0.0540	0.010	5.160	0.000
answer_13().day_Sat	0.0355	0.010	3.590	0.000

```

=====
Omnibus:                214.408      Durbin-Watson:
Prob(Omnibus):          0.000      Jarque-Bera (JB):
Skew:                   3.189      Prob(JB):
Kurtosis:               28.623      Cond. No.
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correct.
[2] The smallest eigenvalue is 2.97e-29. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.
"""

```

1.3.3 Question 15

Fit a model with only the size as an explanatory variable

```

In [32]: def answer_15():
          #write your code here
          data_tips=answer_13()
          model=sm.ols(formula="data_tips.tip_rate~data_tips['size']",data=data_tips)
          return model.summary()

answer_15()

Out[32]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

OLS Regression Results

```

=====
Dep. Variable:      data_tips.tip_rate      R-squared:      0.
Model:              OLS                    Adj. R-squared:  0.
Method:             Least Squares          F-statistic:    5.
Date:               Tue, 18 Jul 2017        Prob (F-statistic): 0.
Time:               13:47:08               Log-Likelihood:  33.
No. Observations:   244                    AIC:           -6.
Df Residuals:       242                    BIC:           -6.
Df Model:           1
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[95.0%
Intercept	0.1844	0.011	16.475	0.000	0.16
data_tips['size']	-0.0092	0.004	-2.245	0.026	-0.01

```

=====
Omnibus:            220.122      Durbin-Watson:      2.
Prob(Omnibus):      0.000      Jarque-Bera (JB):    7618.
Skew:               3.308      Prob(JB):           0.
Kurtosis:           29.562      Cond. No.           1.
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correct
"""

```

1.3.4 Question 16

Use a stepwise algorithm with the AIC statistic as a variable selection process to select a good model. Start from the full model of question 13. What do you remark?

```

In [34]: # this function is provided since there is no specific method for aicstat
# the source code was taken from : http://planspace.org/20150423-forward_s
import statsmodels.formula.api as smf

def forward_selected(data, response):
    """Linear model designed by forward selection.

    Parameters:
    -----
    data : pandas DataFrame with all possible predictors and response

    response: string, name of response column in data

    Returns:
    -----
    model: an "optimal" fitted statsmodels linear model
           with an intercept
    """

```

```

selected by forward selection
evaluated by adjusted R-squared
"""
remaining = set(data.columns)
remaining.remove(response)
selected = []
current_score, best_new_score = 0.0, 0.0
while remaining and current_score == best_new_score:
    scores_with_candidates = []
    for candidate in remaining:
        formula = "{} ~ {}".format(response,
                                     ' + '.join(selected + [candidate]))
        score = smf.ols(formula, data).fit().rsquared_adj
        scores_with_candidates.append((score, candidate))
    scores_with_candidates.sort()
    best_new_score, best_candidate = scores_with_candidates.pop()
    if current_score < best_new_score:
        remaining.remove(best_candidate)
        selected.append(best_candidate)
        current_score = best_new_score
formula = "{} ~ {}".format(response,
                            ' + '.join(selected))
model = smf.ols(formula, data).fit()
return model

```

```

In [35]: def answer_16():
        data=answer_13().iloc[:,2:10]
        model = forward_selected(data, 'tip_rate')
        return model.model.formula ,model.rsquared_adj
answer_16()

```

```

Out[35]: ('tip_rate ~ size + day_Sat ', 0.022243088905636199)

```

Home work Explore the notion of interaction between the Gender and the smoking habit by including explicitly this interaction into a model with size, sex, smoke

1.3.5 Question 17

Check the linear relationship between the tip and the total bill, seen at question 7, with a linear model and interpret the quality of this model

```

In [37]: def answer_17():
        # write your code here
        model=sm.ols(formula="data_tips.tip_rate~data_tips.total_bill",data=data_tips)
        return model.summary()
answer_17()

```

```

Out[37]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
Dep. Variable:          data_tips.tip_rate      R-squared:          0.999
Model:                  OLS                    Adj. R-squared:      0.999
Method:                 Least Squares          F-statistic:         35.85
Date:                  Tue, 18 Jul 2017         Prob (F-statistic):  5.85e-07
Time:                  13:54:37                Log-Likelihood:      35.85
No. Observations:      244                    AIC:                 -6.00
Df Residuals:          242                    BIC:                 -6.00
Df Model:              1
Covariance Type:       nonrobust
=====
                                coef      std err          t      P>|t|      [95.0%
-----+-----
Intercept              0.2068        0.009     22.975     0.000      0.198
data_tips.total_bill   -0.0023        0.000    -5.599     0.000     -0.003
=====
Omnibus:              224.802    Durbin-Watson:      2.000
Prob(Omnibus):        0.000    Jarque-Bera (JB):   7863.000
Skew:                 3.420    Prob(JB):           0.000
Kurtosis:             29.957    Cond. No.           1.00e+01
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correct
"""

```

```
In [ ]:
```